
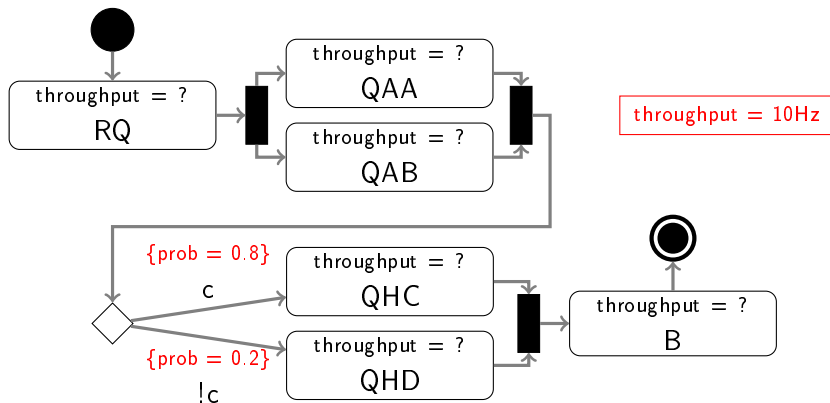


Running example: travel search engine



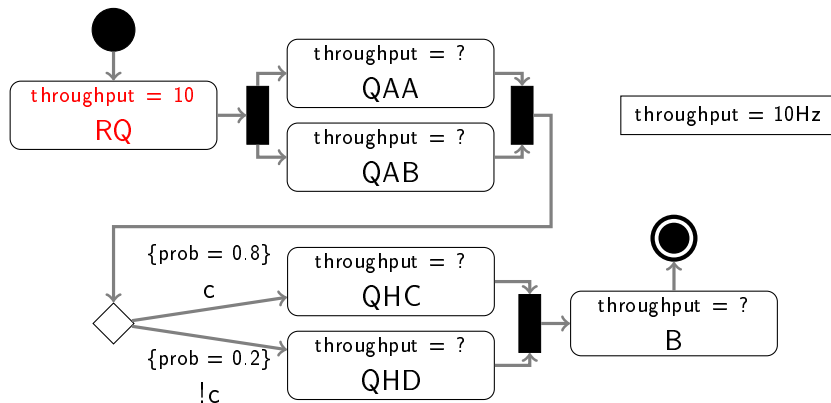
Throughput inference

We add a global requirement and annotate decision branches with their local probabilities.



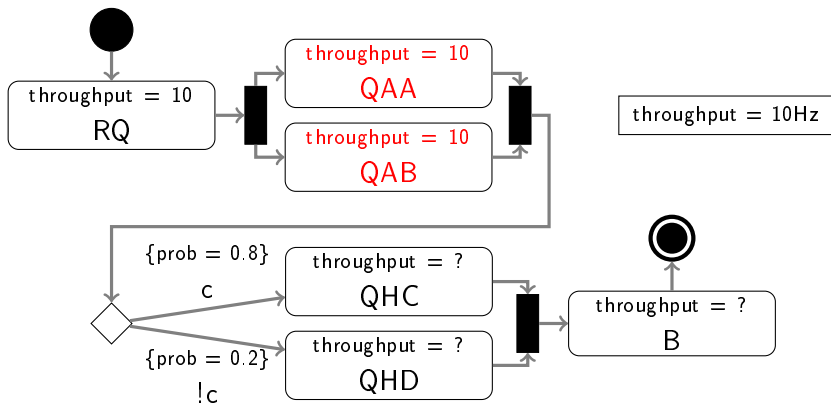
Throughput inference

Requests flow sequentially from the initial node.



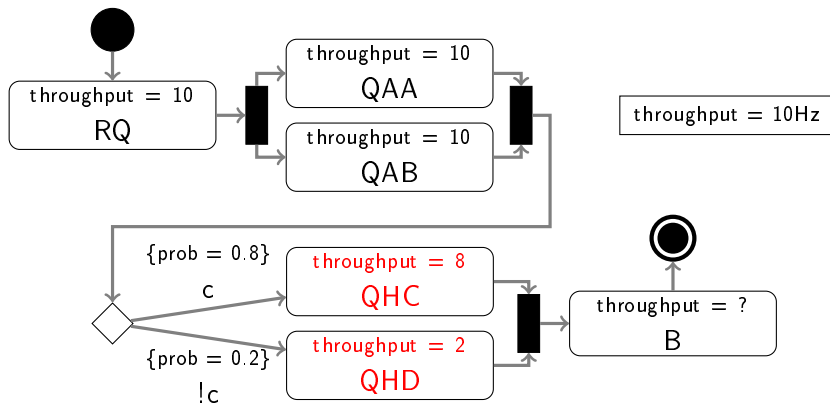
Throughput inference

Forks clone the same request into several concurrent ones, and joins wait for all incoming paths.



Throughput inference

Decisions send the request to a branch, under some probability.



Throughput inference

Merges bring together requests from different branches: we're done!

